




# DM2112 DIGITAL ENTERTAINMENT SYSTEMS




## Operating System Concepts Pt 1

# Today's Menu

- What is an Operating System?
- Interrupts
- Process Control Block
- Process Scheduling
- CPU Scheduling



# WHAT IS AN OPERATING SYSTEM?



# Operating Systems

- What's an Operating System (also called "OS")?
- It's a computer program that:
  - Manages all resources.
  - Decides between conflicting requests for efficient and fair resource use.
  - Controls execution of programs to prevent errors and improper use of the computer.
- A program that acts as an intermediary between a user of a computer and the computer hardware.



# Operating Systems

What it does for you:

- Executes your programs.
- Make a system more convenient for you to use.
- Makes solving problems easier.
- An operating system is a **resource allocator**.
  - Manages all resources
  - Decides between conflicting requests for efficient and fair resource use
- An operating system is a **control program**.
  - Controls execution of programs to prevent errors and improper use of the computer

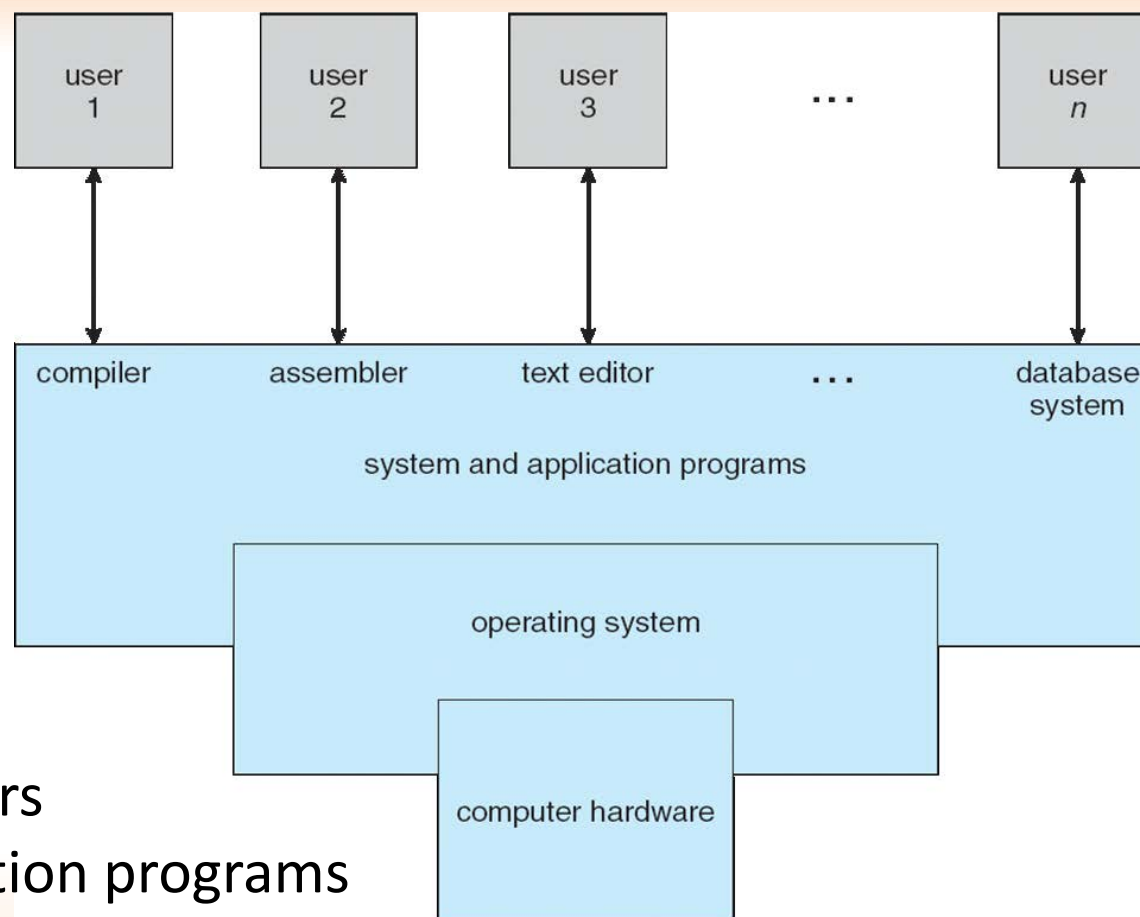


# Bootstrap / Computer Startup

- How is the operating system loaded?
  - **bootstrap program** is loaded at power-up or reboot (so that's why we say we “boot” up a computer).
    - Stored in ROM or EPROM
    - Loads basic part of OS (i.e. kernel) and starts execution



# The Components of a Computer System



1. The users
2. Application programs
3. Operating System
4. Actual hardware



# The Components of a Computer System

A computer system can be divided into four components:

➤ **Hardware**

- provides basic computing resources
- CPU, memory, I/O devices

➤ **Operating system**

- Controls and coordinates use of hardware among various applications and users

➤ **Application programs**

- Define the ways in which the system resources are used to solve the computing problems of the users
- Word processors, compilers, web browsers, database systems, video games

➤ **Users**


- People, machines, other computers

- The operating system handles the interaction between the hardware, the application programs, and the users.






# INTERRUPTS



# Interrupts



- A computer system has many devices competing for the CPU's attention.
  - Each kind of device has its own device controller
  - Each device controller has its own buffer
  - The CPU moves data to/from main memory to the device controller's buffers, and the controller transfers the data to/from the appropriate device
  - The device controller informs the CPU that it has finished its I/O operation by causing an interrupt




# What Happens When an Interrupt Occurs?

- The address of the interrupted instruction must be saved.
  - Incoming interrupts are *disabled* while another interrupt is being processed to prevent a *lost interrupt*
- The operating system preserves the state of the CPU by storing the register values (this CPU state is called the “context”).
- The operating system determines which type of interrupt has occurred.
  - Polling (i.e. software interrupt)
  - Vectored interrupt (i.e. hardware interrupt)
- Separate segments of code determine what action should be taken for each type of interrupt.
- An operating system is **interrupt driven**.




# Hardware Interrupts (IRQ)

- An interrupt request (IRQ) is a common way that a piece of hardware informs the OS that an important event has occurred.
  - Versus having CPU continuously poll ('ask') each device if there is an event
    - Polling is a costly process (Why?)
  - Fast response from CPU when triggered



# Analogy – Phone Operator


- Long long ago (in a galaxy far far away?)



- But what if?



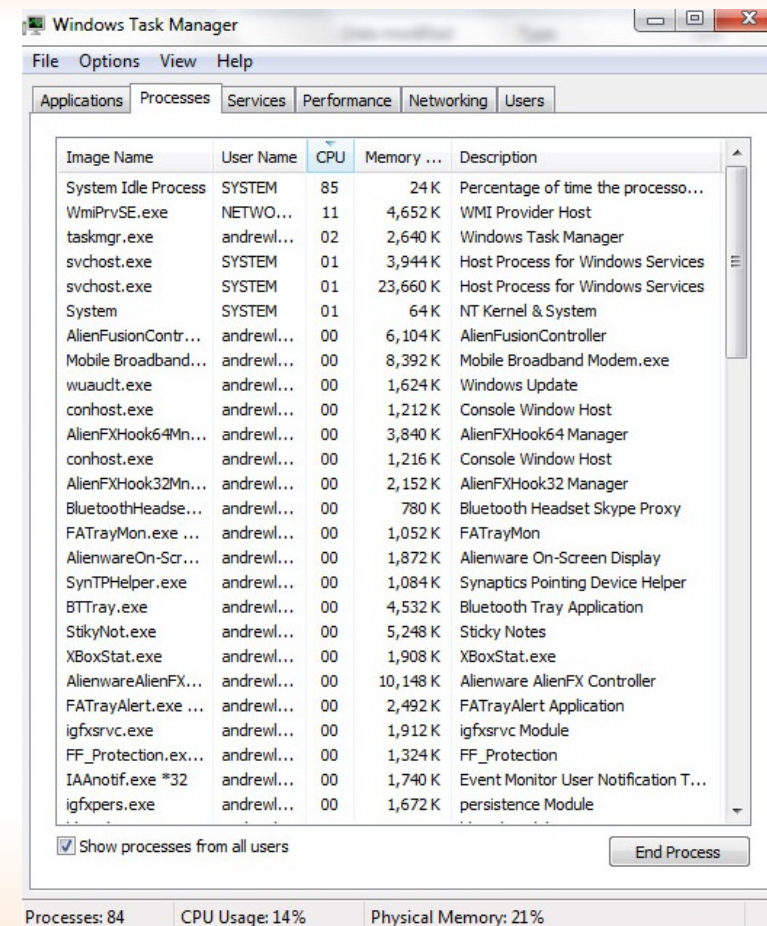
# PROCESSES





# Concept of a Process

- An operating system executes a variety of programs:
  - Batch system – jobs
  - Time-shared systems – user programs or tasks
- We usually use the terms *job* and *process* almost interchangeably.



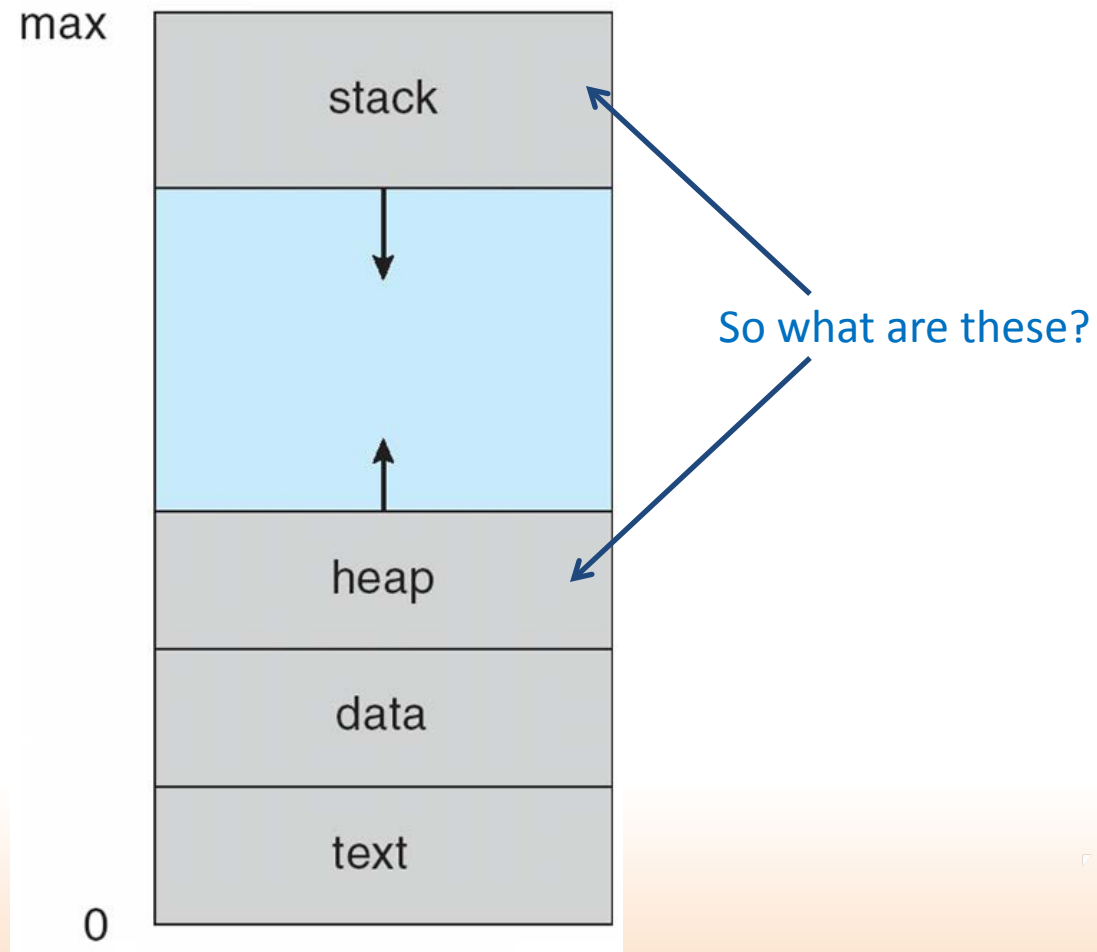
# Concept of a Process

- Process – a program in execution; process execution must progress in sequential fashion.
- A program on its own is just a bunch of code. It is only called a process when a program is running (i.e. when it is executing).
- A process includes:
  - A program counter
  - A stack
  - A data section



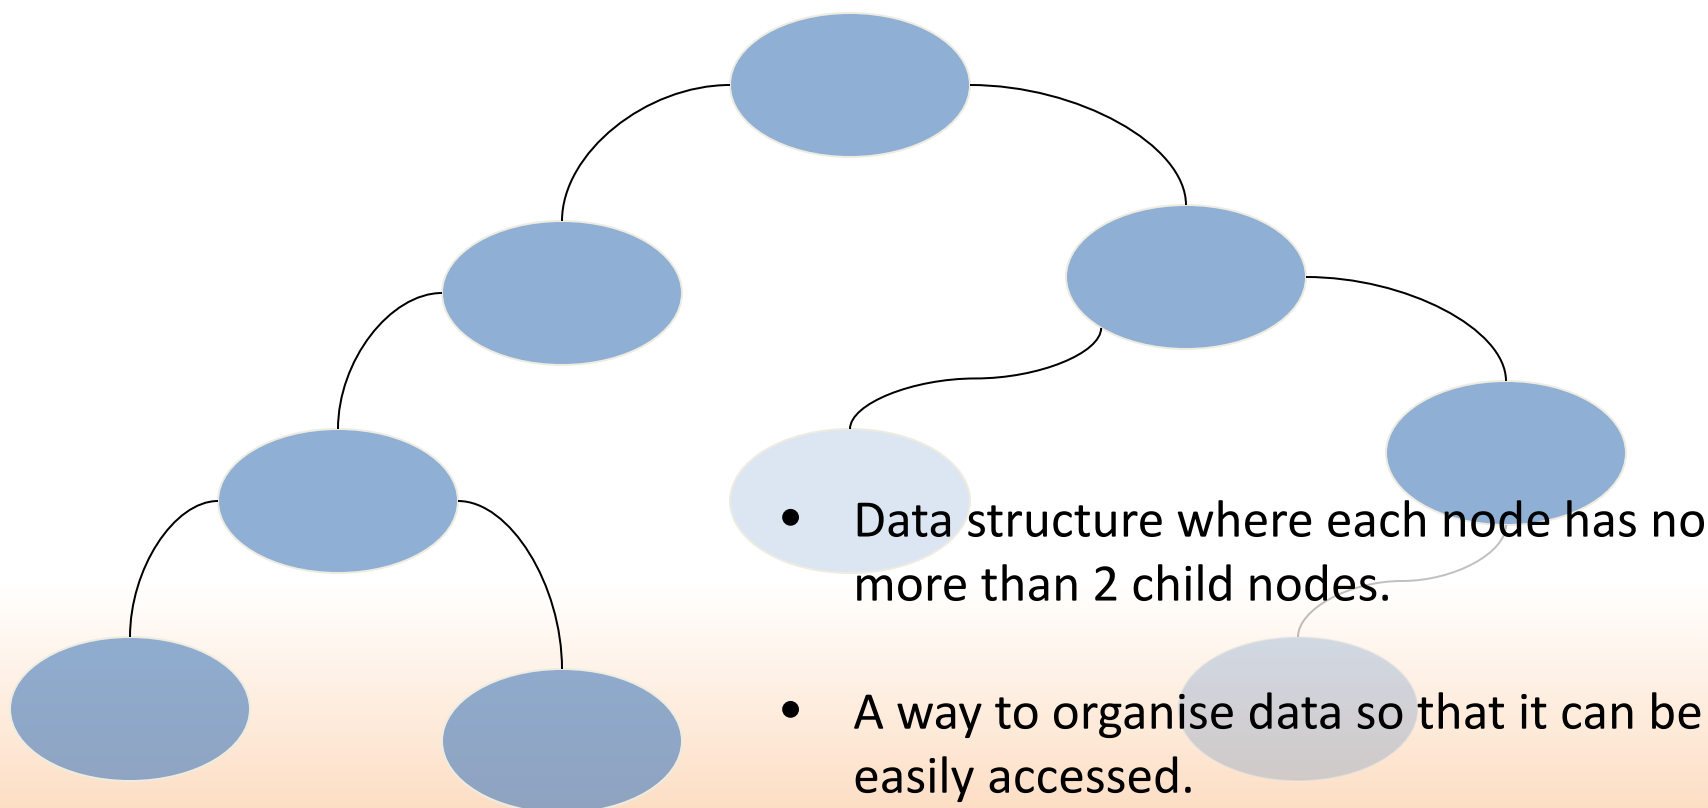


# A Process in Memory




# Heap

- A form of binary tree.
- A heap is used to store long-term data.



# Stack

- A reversed queue system.
  - Last in, First Out (LIFO or First In, Last Out - FILO)
- A stack is used to store short-term data.



Like a stack of paper

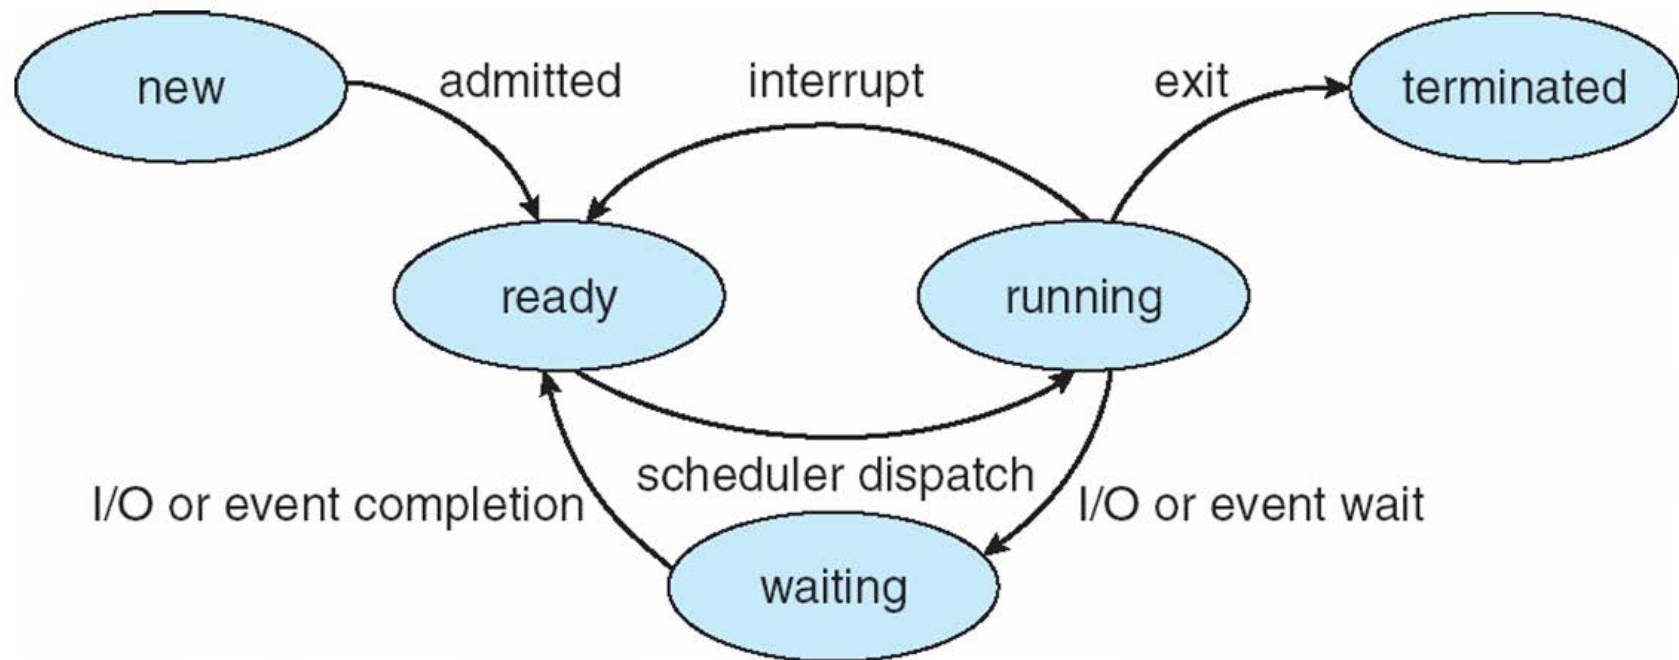


# Process States

- As a process executes, it changes its *state*.
  - **New:** The process is being created
  - **Running:** Instructions are being executed
  - **Waiting:** The process is waiting for some event to occur
  - **Ready:** The process is waiting to be assigned to a processor
  - **Terminated:** The process has finished execution

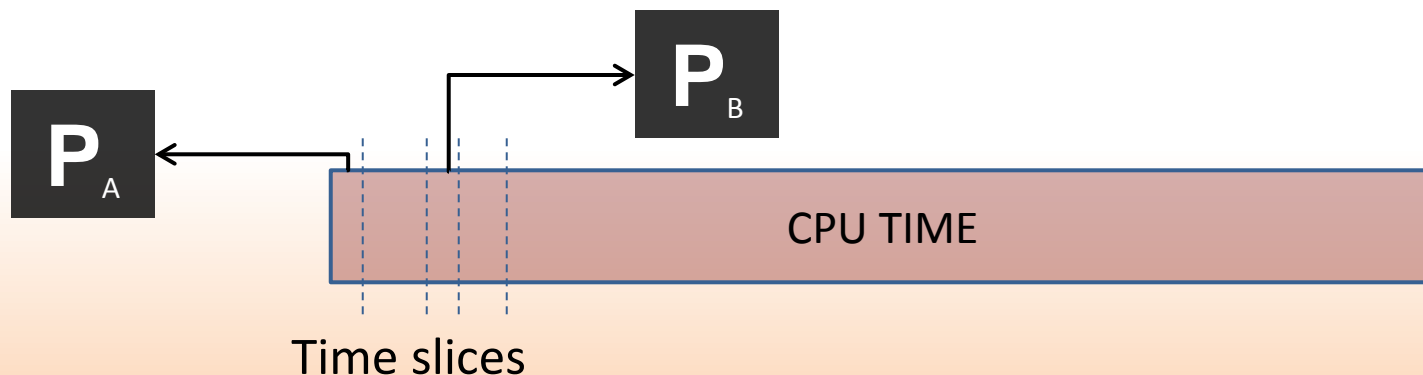


# Diagram of Process State



# Process State

- Only 1 process can be running on a processor (i.e. CPU) at any given time.
- Multitasking is achieved by moving the current process out of the CPU, and then letting another process execute.
- CPU switches between processes every few milliseconds.
  - From user's point of view, all the programs are running at the same time
- The switching of processes every few seconds is called time slicing



# Process Creation

- Processes are arranged in a tree-like structure.
  - Can be parent or child of other processes
  - Root process has no parent (i.e. a root process is a kernel process)
- Process identified and managed via **process identifier (pid)**.
- Execution.
  - Parent and children execute concurrently
  - Parent waits until children terminate



# Example of Processes Running in a System

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
227	root	20	0	0	0	0	S	0	0.0	1:07.17	usb-storage
30904	root	20	0	19352	1344	960	R	0	0.0	0:00.05	top
1	root	20	0	24136	2300	1332	S	0	0.1	0:00.70	init
2	root	20	0	0	0	0	S	0	0.0	0:00.18	kthreadd
3	root	20	0	0	0	0	S	0	0.0	0:42.98	ksoftirqd/0
6	root	RT	0	0	0	0	S	0	0.0	0:00.00	migration/0
7	root	RT	0	0	0	0	S	0	0.0	0:00.00	migration/1
9	root	20	0	0	0	0	S	0	0.0	0:15.45	ksoftirqd/1
11	root	0	-20	0	0	0	S	0	0.0	0:00.00	cpuset
12	root	0	-20	0	0	0	S	0	0.0	0:00.00	khelper
13	root	0	-20	0	0	0	S	0	0.0	0:00.00	netns
14	root	20	0	0	0	0	S	0	0.0	0:03.89	kworker/u:1
15	root	20	0	0	0	0	S	0	0.0	0:01.13	sync_supers
16	root	20	0	0	0	0	S	0	0.0	0:00.07	bdi-default
17	root	0	-20	0	0	0	S	0	0.0	0:00.00	kintegrityd
18	root	0	-20	0	0	0	S	0	0.0	0:00.00	kblockd
19	root	0	-20	0	0	0	S	0	0.0	0:00.00	kacpid
20	root	0	-20	0	0	0	S	0	0.0	0:00.00	kacpi_notify






# Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
  - Process' resources are deallocated by operating system (the memory that it was using is cleared)
- Parent may terminate execution of children processes (**abort**) for the following reasons:
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent process is exiting



# PROCESS CONTROL BLOCK



# Process Control Block (PCB)


Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

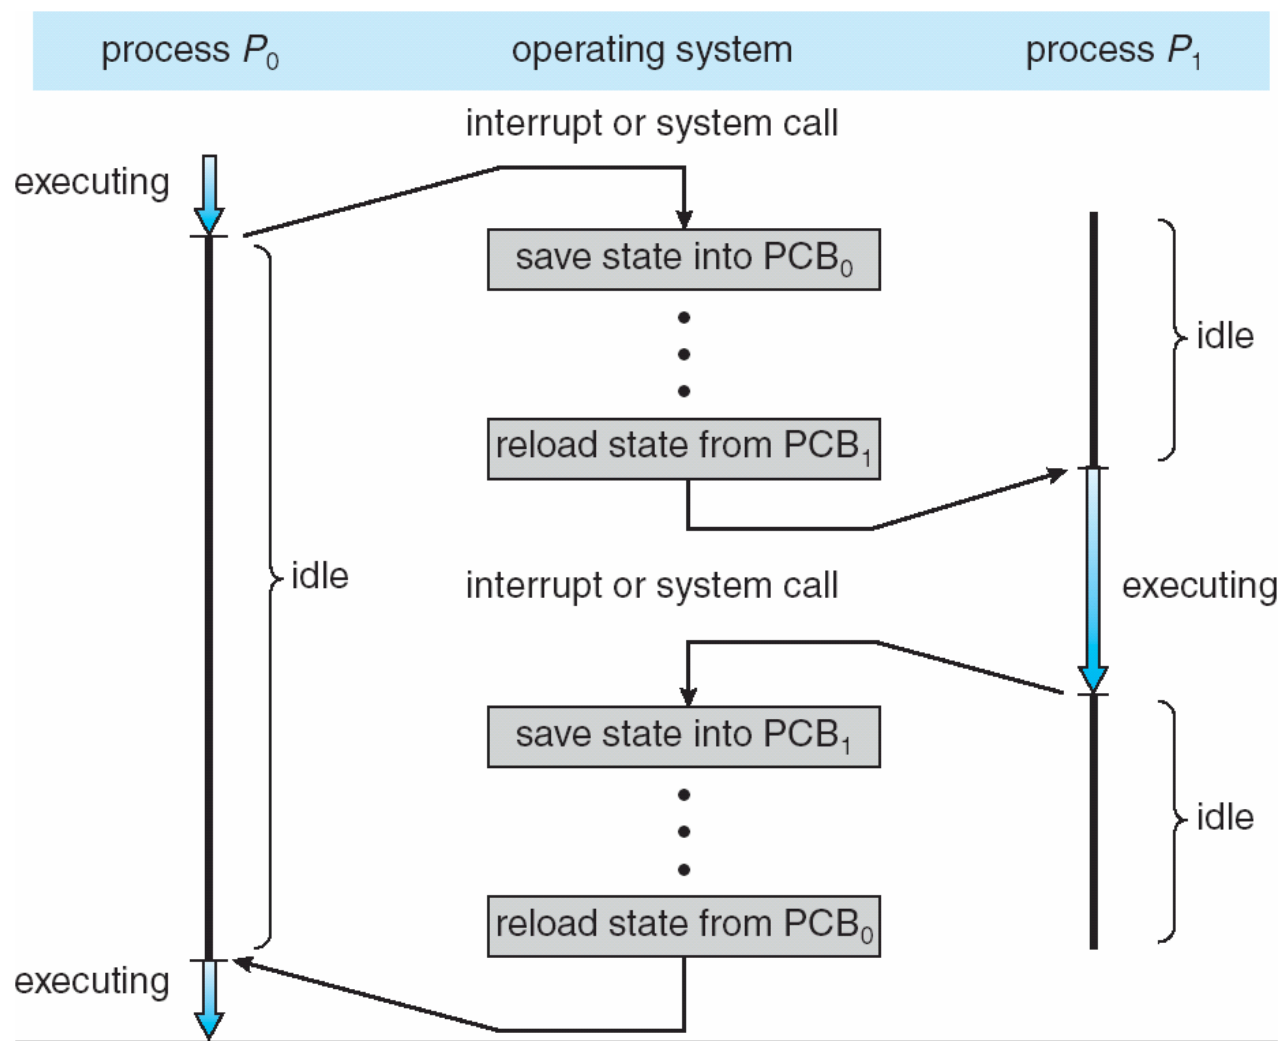
# Process Control Block (PCB)




# PROCESS SCHEDULING



# CPU Switches From Process to Process



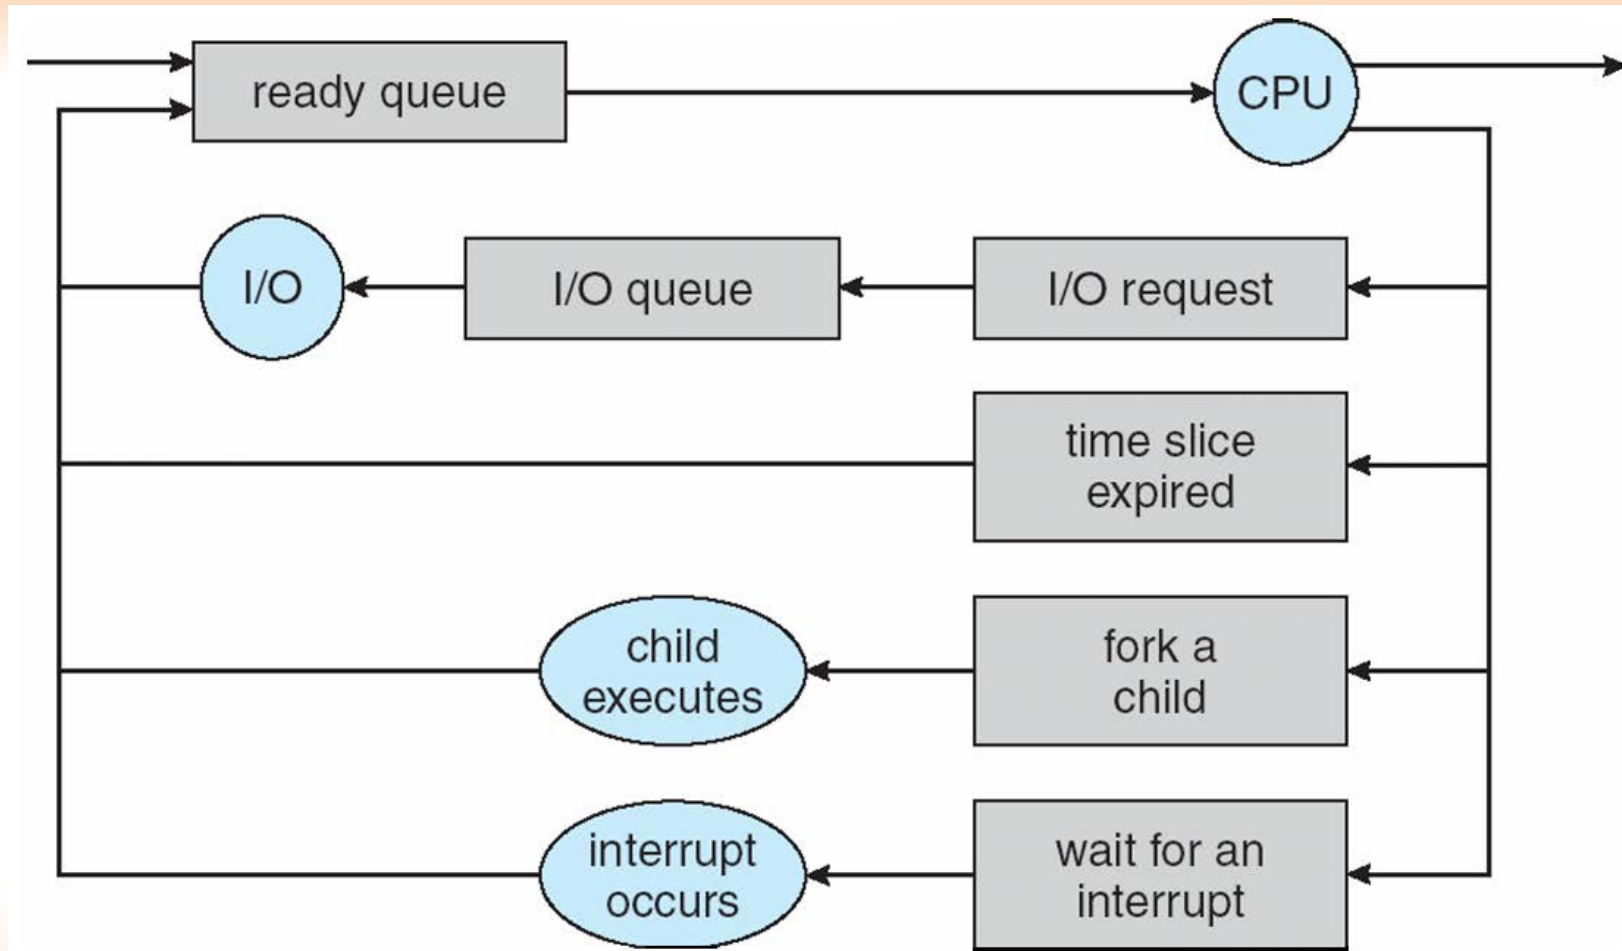
# Process Scheduling Queues



Processes move among  
the various queues

- **Job queue** – set of all processes in the system.
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute.
- **Device queues** – set of processes waiting for an I/O device.

# Representation of Process Scheduling





# Schedulers

- **Long-term scheduler**
  - Job scheduler
  - Selects which processes should be brought into the ready queue
- **Short-term scheduler**
  - CPU scheduler
  - Selects which process should be executed next and allocates CPU



# Schedulers

- **Short-term scheduler** is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast)
- **Long-term scheduler** is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)



# Processes & Scheduling

- Processes can be described as either:
  - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- CPU burst is the amount of time a process uses the processor before the process is no longer ready.
  - e.g. until the time when the process needs to wait for Input/Output such as waiting to write something to the harddisk, or waiting for user to type something on the keyboard
- I/O burst is the time a process is accessing the I/O devices until it has to wait for the CPU to be ready again.
  - e.g. while the process is accessing the I/O the CPU attends to other processes




# Context Switch

- When CPU switches to another process:
  - System must save the state of the old process
  - Load the saved state for the new process via a **context switch**
- Context of a process is represented in the PCB.
- Context switch time is overhead; the system does no useful work while switching.




# CPU SCHEDULING

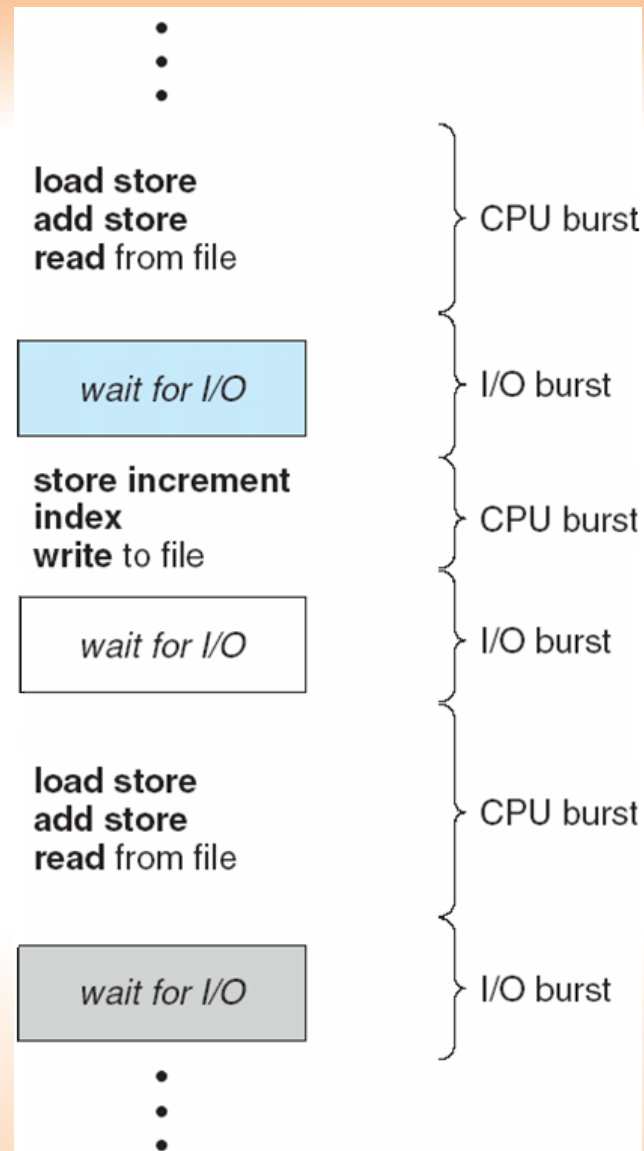


# Basic Concepts

- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait.



# Alternating Sequence of CPU And I/O Bursts



# The Concept of Scheduling



$A \rightarrow B \rightarrow C$  or

$B \rightarrow A \rightarrow C$  or

$B \rightarrow C \rightarrow A$  or

$C \rightarrow B \rightarrow A$

- Which job should we do first?
- Which job should come next?
- We need to consider things such as:
  - Which job came first?
  - Which job takes the shortest time to complete?





# Scheduling Criteria

- **CPU utilization** – aim is to keep the CPU as busy as possible.
- **Throughput** – Number of processes that complete their execution per time unit.
- **Turnaround time** – amount of time to execute a particular process.
- **Waiting time** – amount of time a process has been waiting in the ready queue.
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment).



# Scheduling Algorithm Optimization Criteria


- Maximum CPU utilization
- Maximum throughput
- Minimum turnaround time
- Minimum waiting time
- Minimum response time



# First-Come, First-Served Scheduling (FCFS)

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$
- The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$




# FCFS Scheduling (Continued)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- So now the Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case.



# Shortest-Job-First (SJF) Scheduling

- Associates with each process the length of its next CPU burst (i.e. the CPU burst time when the process is next using the CPU).
- Uses these lengths to schedule the process with the shortest time.
- SJF is optimal – gives minimum average waiting time for a given set of processes.
  - The difficulty is knowing the length of the next CPU request




# Example of SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	6
$P_2$	2.0	8
$P_3$	4.0	7
$P_4$	5.0	3

Note that  $P_1$  runs first,  
Because it arrived at time 0.0 where no other processes have arrived

SJF scheduling chart:



By time 5.0, all other jobs have arrived

$$\text{Average waiting time} = ((0-0) + (6-5) + (9-4) + (16-2)) / 4 = (0+1+5+14)/4 = 5$$



# FYI: Other Types of Scheduling

- Priority
  - Round Robin
  - Lottery
  - Least Slack Time
- 
- Not discussed in this module but do read up on them for your own enrichment.



# Acknowledgements

- Slides adapted from Operating System Concepts (8<sup>th</sup> Edition), by Abraham Silberschatz, Greg Gagne, and Peter Baer Galvin. John Wiley & Sons Inc., ISBN 0-470-12872-0





# More Next Time

- Questions?

